

Fake News detection

Naman Agarwal

Overview

To use various natural language processing techniques and machine learning algorithms to classify fake news articles using sci-kit libraries from python.

Goals

1. Binary classification task (true, false)
2. Six-way classification task (pants on fire, false, mostly false, half-true, mostly true, true)

Prerequisites

1. Python 3.6
2. You will also need to download and install below packages after you install either python or anaconda from the steps above
 - Sklearn (scikit-learn)
 - numpy
 - Scipy
 - NLTK
 - Seaborn
 - Gensim
 - Matplotlib
 - Pandas

File Description

DataPrep.py

This file contains all the pre processing functions needed to process all input documents and texts. First we read the train, test and validation data files then performed some preprocessing like tokenizing, stemming etc. There are some exploratory data analysis is performed like response variable distribution and data quality checks like null or missing values etc.

FeatureSelection.py

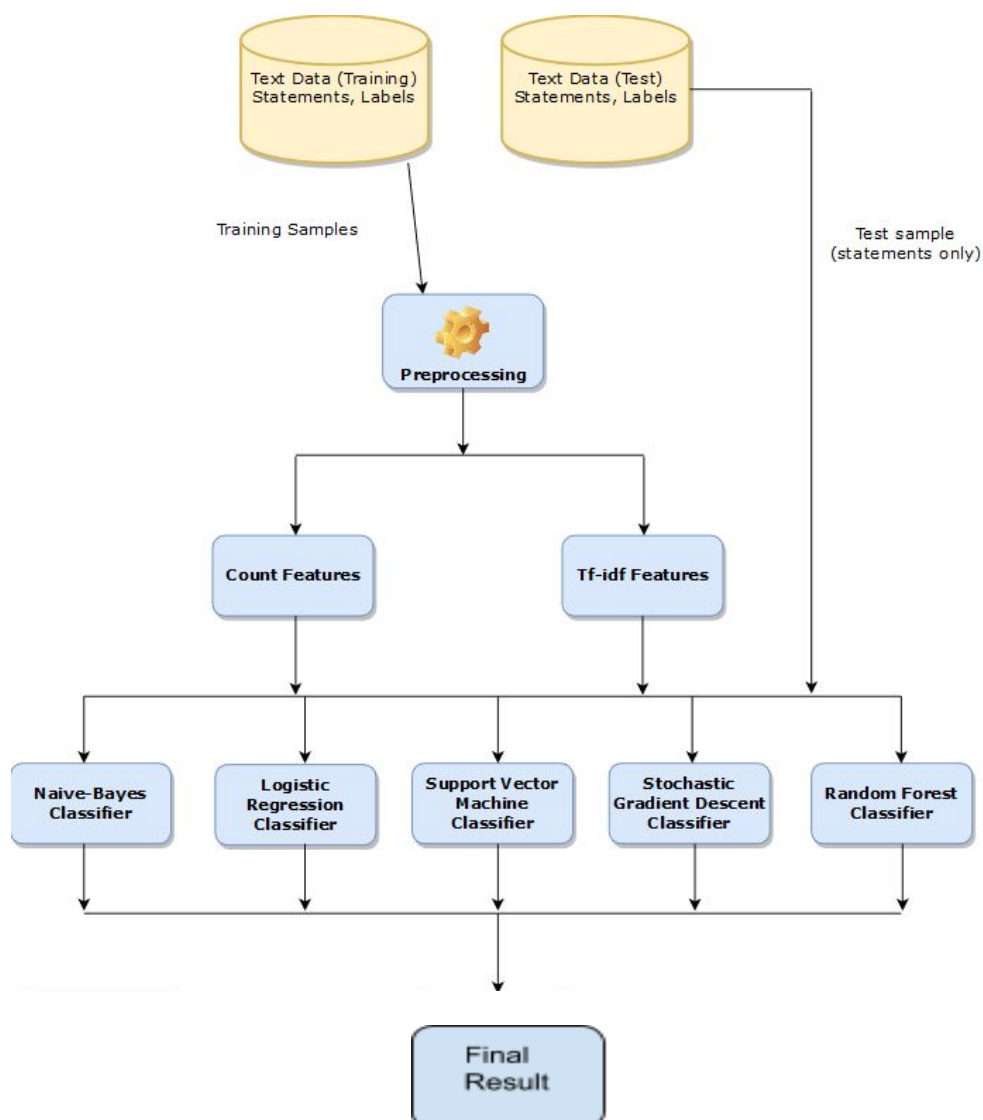
This file is used to perform feature extraction and selection methods from sci-kit learn python libraries. For feature selection, methods like simple bag-of-words and n-grams and then term frequency like tf-idf weighting are used.

classifier.py

Here I have built all the classifiers for predicting the fake news detection. The extracted features are fed into different classifiers. I have used Naive-bayes, Logistic Regression, Linear SVM, Stochastic gradient descent and Random Forest classifiers from sklearn. Each of the extracted features were used in all of the classifiers. Once fitting the model,

we compared the f1 score and checked the confusion matrix. After fitting all the classifiers, 2 best performing models were selected as candidate models for fake news classification. We have performed parameter tuning by implementing GridSearchCV method on these candidate models and chosen best performing parameters for these classifier. Finally selected model was used for fake news detection with the probability of truth. In Addition to this, We have also extracted the top 50 features from our term-frequency tfidf vectorizer to see what words are most important in each of the classes. We have also used Precision-Recall and learning curves to see how training and test set performs when we increase the amount of data in our classifiers.

Process Flow Chart



Performance

Below are the confusion matrices and F1 score for binary classifier using:

1. n-grams & tfidf

Naive bayes

[841 3647]

[427 5325]

f1-Score: 0.723262051071

Logistic regression

[1617 2871]

[1097 4655]

f1-Score: 0.70113000531

SVM

[2016 2472]

[1524 4228]


f1-Score: 0.67909201429

Sgdclassifier

[10 4478]

[13 5739]

f1-Score: 0.718731637053



random forest

[1979 2509]

[1630 4122]

f1-Score: 0.665720333284

2. Bag of words

Naive bayes

[2118 2370]

[1664 4088]

f1-Score: 0.669611539651

Logistic regression

[2252 2236]

[1933 3819]

f1-Score: 0.646909097798

SVM

[2260 2228]

[2246 3506]

f1-score: 0.610468748792

Sgdclassifier

[2414 2074]

[2042 3710]

f1-Score: 0.640874558778

random forest classifier

[1821 2667]

[1192 4560]

f1-Score: 0.702651511011

Below is the classification report for binary classifier:

classifier_binary.py x classifier_multiclass.py x untitled						
		precision	recall	f1-score	support	
190						
191						
192						
193	False	0.72	0.19	0.30	1169	
194	True	0.58	0.94	0.71	1382	
195						
196	accuracy			0.59	2551	
197	macro avg	0.65	0.56	0.51	2551	
198	weighted avg	0.64	0.59	0.52	2551	
199						

Naive Bayes

		precision	recall	f1-score	support	
	False	0.64	0.39	0.49	1169	
	True	0.61	0.81	0.70	1382	
	accuracy			0.62	2551	
	macro avg	0.62	0.60	0.59	2551	
	weighted avg	0.62	0.62	0.60	2551	

Logistic Regression

		precision	recall	f1-score	support	
209						
210						
211	False	0.61	0.47	0.53	1169	
212	True	0.62	0.74	0.68	1382	
213						
214	accuracy			0.62	2551	
215	macro avg	0.61	0.61	0.60	2551	
216	weighted avg	0.62	0.62	0.61	2551	
217						

SVM Classifier

217						
218			precision	recall	f1-score	support
219						
220		False	0.00	0.00	0.00	1169
221		True	0.54	1.00	0.70	1382
222						
223		accuracy			0.54	2551
224		macro avg	0.27	0.50	0.35	2551
225		weighted avg	0.29	0.54	0.38	2551
226						

SGD Classifier

226						
227			precision	recall	f1-score	support
228						
229		False	0.58	0.47	0.52	1169
230		True	0.61	0.71	0.66	1382
231						
232		accuracy			0.60	2551
233		macro avg	0.59	0.59	0.59	2551
234		weighted avg	0.60	0.60	0.59	2551
235						

Random Forest

Below are the confusion matrices and F1 score for multi class classifier using:

1. Bag of words

Naive bayes

```
[[255 377 538 351 13 120]
 [275 519 582 414 15 190]
 [251 366 713 561 10 213]
 [171 294 615 593 9 280]
 [144 246 218 137 35 59]
 [115 281 456 550 18 256]]
```

f1-Score: [0.17810937 0.25459002 0.27235932 0.25946755 0.07458623
0.18311458]

Logistic regression

```
[[314 375 380 325 70 190]
 [295 586 403 332 108 271]
 [318 406 562 459 69 300]
 [251 349 449 483 52 378]
 [145 228 148 113 114 91]]
```



[169 324 344 443 53 343]]

f1-Score: [0.19963917 0.27494763 0.25496941 0.2346005 0.17429044
0.21102356]

SVM

[[316 356 350 295 106 231]

[307 501 382 354 146 305]

[340 388 506 439 103 338]

[281 365 429 440 81 366]

[143 193 141 102 144 116]

[207 302 323 403 87 354]]

f1-Score: [0.19441198 0.24408367 0.23795384 0.22010213 0.1911165
0.20903274]

Sgdclassifier

[[251 379 413 299 128 184]

[251 593 441 300 155 255]

[245 416 627 450 118 258]

[223 332 551 439 92 325]

[117 212 168 104 154 84]

[172 339 391 380 83 311]]

f1-Score: [0.1685202 0.27561171 0.26085534 0.21971733 0.1951177
0.19916655]


Random Forest

[[182 635 378 331 5 123]

[150 913 431 330 20 151]

[153 651 627 486 7 190]

[102 549 446 611 7 247]



```
[ 68 401 159 116 56 39]
[ 80 490 356 497 3 250]]
```

```
f1-Score: [0.15225759 0.32428993 0.27751728 0.28154819 0.11931054
0.1866272 ]
```

2. N-grams and tfidf

Naive bayes

```
[[ 60 398 884 292 0 20]
 [ 48 577 1022 315 0 33]
 [ 52 400 1169 456 0 37]
 [ 17 330 1056 514 0 45]
 [ 21 281 406 122 2 7]
 [ 19 315 835 454 0 53]]
```

```
f1-Score: [0.06417057 0.2684663 0.31294388 0.24858862 0.00440801
0.05670925]
```


Logistic regression

```
[[201 438 557 375 0 83]
 [190 653 646 385 2 119]
 [177 437 771 574 2 153]
 [116 364 672 618 2 190]
 [107 297 242 145 14 34]
 [ 93 327 499 545 1 211]]
```

```
f1-Score: [0.15806962 0.28980254 0.28017956 0.26830834 0.03294609
0.17093725]
```

SVM

```
[[311 334 424 351 54 180]
 [310 476 497 387 83 242]
```



[276 340 656 532 45 265]
[210 295 501 565 34 357]
[155 190 173 131 119 71]
[177 280 365 483 35 336]]

f1-Score: [0.20104128 0.24343244 0.27733517 0.25611377 0.19658431
0.21453804]

Sgdclassifier

[[326 319 439 335 51 184]
[324 458 501 371 85 256]
[289 332 654 515 45 279]
[217 285 515 550 38 357]
[164 194 162 126 116 77]
[193 267 375 460 34 347]]

f1-Score: [0.20573008 0.23811802 0.27477963 0.25458604 0.19148554
0.21831309]

Random Forest

[[146 666 282 366 41 153]
[152 902 309 415 53 164]
[156 732 355 639 33 199]
[115 602 310 664 25 246]
[80 387 115 129 69 59]
[101 560 205 521 23 266]]

f1-Score: [0.12007153 0.30902868 0.19245476 0.28280816 0.1261867
0.19239621]

Below is the classification report for multiclass classifier:

		precision	recall	f1-score	support
barely-true		0.24	0.03	0.06	212
false		0.28	0.31	0.29	249
half-true		0.23	0.58	0.33	265
mostly-true		0.23	0.25	0.24	241
pants-fire		0.00	0.00	0.00	92
true		0.12	0.01	0.03	208
accuracy				0.24	1267
macro avg		0.18	0.20	0.16	1267
weighted avg		0.21	0.24	0.19	1267

Naive Bayes

		precision	recall	f1-score	support
barely-true		0.27	0.14	0.18	212
false		0.27	0.34	0.30	249
half-true		0.23	0.35	0.28	265
mostly-true		0.22	0.30	0.25	241
pants-fire		0.50	0.01	0.02	92
true		0.25	0.14	0.18	208
accuracy				0.24	1267
macro avg		0.29	0.21	0.20	1267
weighted avg		0.27	0.24	0.23	1267

Logistic Regression

		precision	recall	f1-score	support
barely-true		0.25	0.21	0.23	212
false		0.29	0.28	0.29	249
half-true		0.26	0.33	0.29	265
mostly-true		0.22	0.27	0.25	241
pants-fire		0.15	0.07	0.09	92
true		0.24	0.22	0.23	208
accuracy				0.25	1267
macro avg		0.24	0.23	0.23	1267
weighted avg		0.25	0.25	0.25	1267

SVM Classifier


		precision	recall	f1-score	support
158					
159					
160	barely-true	0.27	0.23	0.25	212
161	false	0.33	0.28	0.30	249
162	half-true	0.24	0.31	0.27	265
163	mostly-true	0.22	0.28	0.25	241
164	pants-fire	0.17	0.09	0.12	92
165	true	0.23	0.20	0.21	208
166					
167	accuracy			0.25	1267
168	macro avg	0.24	0.23	0.23	1267
169	weighted avg	0.25	0.25	0.25	1267
170					

SGD Classifier

		precision	recall	f1-score	support
171					
172					
173	barely-true	0.30	0.12	0.17	212
174	false	0.24	0.46	0.32	249
175	half-true	0.23	0.17	0.19	265
176	mostly-true	0.23	0.35	0.28	241
177	pants-fire	0.29	0.08	0.12	92
178	true	0.23	0.14	0.18	208
179					
180	accuracy			0.24	1267
181	macro avg	0.25	0.22	0.21	1267
182	weighted avg	0.25	0.24	0.22	1267
183					

Random Forest

The best classifiers were Logistic Regression and Random Forest and the learning curves for these classifiers are as follows:



References

1. Jones E, Oliphant T, Peterson P, others. SciPy: Open source scientific tools for Python [Internet]. 2001. Available from: "<http://www.scipy.org/>"
2. Oliphant TE. A guide to NumPy. Vol. 1. Trelgol Publishing USA; 2006.
3. McKinney W, others. Data structures for statistical computing in python. In: Proceedings of the 9th Python in Science Conference. 2010. p. 51–6.
4. Van Rossum G, Drake Jr FL. Python tutorial. Centrum voor Wiskunde en Informatica Amsterdam, The Netherlands; 1995.
5. Pedregosa F, Varoquaux, Ga"el, Gramfort A, Michel V, Thirion B, Grisel O, et al. Scikit-learn: Machine learning in Python. Journal of machine learning research. 2011;12(Oct):2825–30.
6. Bird, Steven, Edward Loper and Ewan Klein (2009).Natural Language Processing with Python. O'Reilly Media Inc.
7. Hunter JD. Matplotlib: A 2D graphics environment. Computing in science & engineering. 2007;9(3):90–5.